

In [1]:

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn import metrics
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split
from sklearn.cluster import KMeans
```

In [2]:

```
df=pd.read_csv('sales_data_sample.csv',encoding='unicode_escape')
```


In [3]:

```
df.head()
```

Out[3]:

	ORDERNUMBER	QUANTITYORDERED	PRICEEACH	ORDERLINENUMBER	SALES	ORDERDATE
0	10107	30	95.70	2	2871.00	2/2
1	10121	34	81.35	5	2765.90	5/7/200
2	10134	41	94.74	2	3884.34	7/1/200
3	10145	45	83.26	6	3746.70	8/2
4	10159	49	100.00	14	5205.27	10/1

5 rows × 25 columns



In [4]:

```
df.size
```

Out[4]:

70575

In [5]:

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2823 entries, 0 to 2822
Data columns (total 25 columns):
 #   Column                Non-Null Count  Dtype
---  -
 0   ORDERNUMBER           2823 non-null   int64
 1   QUANTITYORDERED       2823 non-null   int64
 2   PRICEEACH             2823 non-null   float64
 3   ORDERLINENUMBER       2823 non-null   int64
 4   SALES                 2823 non-null   float64
 5   ORDERDATE             2823 non-null   object
 6   STATUS                2823 non-null   object
 7   QTR_ID                2823 non-null   int64
 8   MONTH_ID              2823 non-null   int64
 9   YEAR_ID               2823 non-null   int64
10   PRODUCTLINE           2823 non-null   object
11   MSRP                  2823 non-null   int64
12   PRODUCTCODE           2823 non-null   object
13   CUSTOMERNAME          2823 non-null   object
14   PHONE                 2823 non-null   object
15   ADDRESSLINE1          2823 non-null   object
16   ADDRESSLINE2          302 non-null    object
17   CITY                  2823 non-null   object
18   STATE                 1337 non-null   object
19   POSTALCODE            2747 non-null   object
20   COUNTRY               2823 non-null   object
21   TERRITORY             1749 non-null   object
22   CONTACTLASTNAME       2823 non-null   object
23   CONTACTFIRSTNAME      2823 non-null   object
24   DEALSIZE              2823 non-null   object
dtypes: float64(2), int64(7), object(16)
memory usage: 551.5+ KB
```

In [6]:

```
df.shape
```

Out[6]:

```
(2823, 25)
```

In [7]:

```
df.describe().transpose()
```

Out[7]:

	count	mean	std	min	25%	50%	75%
ORDERNUMBER	2823.0	10258.725115	92.085478	10100.00	10180.00	10262.0	10333.
QUANTITYORDERED	2823.0	35.092809	9.741443	6.00	27.00	35.0	43.
PRICEEACH	2823.0	83.658544	20.174277	26.88	68.86	95.7	100.
ORDERLINENUMBER	2823.0	6.466171	4.225841	1.00	3.00	6.0	9.
SALES	2823.0	3553.889072	1841.865106	482.13	2203.43	3184.8	4508.
QTR_ID	2823.0	2.717676	1.203878	1.00	2.00	3.0	4.
MONTH_ID	2823.0	7.092455	3.656633	1.00	4.00	8.0	11.
YEAR_ID	2823.0	2003.815090	0.699670	2003.00	2003.00	2004.0	2004.
MSRP	2823.0	100.715551	40.187912	33.00	68.00	99.0	124.

In [8]:

```
df.columns
```

Out[8]:

```
Index(['ORDERNUMBER', 'QUANTITYORDERED', 'PRICEEACH', 'ORDERLINENUMBER',
      'SALES', 'ORDERDATE', 'STATUS', 'QTR_ID', 'MONTH_ID', 'YEAR_ID',
      'PRODUCTLINE', 'MSRP', 'PRODUCTCODE', 'CUSTOMERNAME', 'PHONE',
      'ADDRESSLINE1', 'ADDRESSLINE2', 'CITY', 'STATE', 'POSTALCODE',
      'COUNTRY', 'TERRITORY', 'CONTACTLASTNAME', 'CONTACTFIRSTNAME',
      'DEALSIZE'],
      dtype='object')
```

In [9]:

```
df.isnull().sum()
```

Out[9]:

ORDERNUMBER	0
QUANTITYORDERED	0
PRICEEACH	0
ORDERLINENUMBER	0
SALES	0
ORDERDATE	0
STATUS	0
QTR_ID	0
MONTH_ID	0
YEAR_ID	0
PRODUCTLINE	0
MSRP	0
PRODUCTCODE	0
CUSTOMERNAME	0
PHONE	0
ADDRESSLINE1	0
ADDRESSLINE2	2521
CITY	0
STATE	1486
POSTALCODE	76
COUNTRY	0
TERRITORY	1074
CONTACTLASTNAME	0
CONTACTFIRSTNAME	0
DEALSIZE	0
dtype:	int64

In [10]:

```
df.dtypes
```

Out[10]:

```
ORDERNUMBER      int64
QUANTITYORDERED  int64
PRICEEACH        float64
ORDERLINENUMBER  int64
SALES            float64
ORDERDATE        object
STATUS           object
QTR_ID           int64
MONTH_ID         int64
YEAR_ID          int64
PRODUCTLINE      object
MSRP             int64
PRODUCTCODE      object
CUSTOMERNAME     object
PHONE            object
ADDRESSLINE1     object
ADDRESSLINE2     object
CITY             object
STATE            object
POSTALCODE       object
COUNTRY          object
TERRITORY        object
CONTACTLASTNAME  object
CONTACTFIRSTNAME object
DEALSIZE         object
dtype: object
```

In [11]:

```
li=['ORDERNUMBER','STATUS','CUSTOMERNAME','PHONE','ADDRESSLINE1','ADDRESSLINE2','CITY',
'STATE','POSTALCODE','TERRITORY','CONTACTLASTNAME','CONTACTFIRSTNAME']
```

In [12]:

```
df.drop(li,inplace=True,axis=1)
```

In [13]:

```
df.columns
```

Out[13]:

```
Index(['QUANTITYORDERED', 'PRICEEACH', 'ORDERLINENUMBER', 'SALES', 'ORDERD
ATE',
      'QTR_ID', 'MONTH_ID', 'YEAR_ID', 'PRODUCTLINE', 'MSRP', 'PRODUCTCOD
E',
      'COUNTRY', 'DEALSIZE'],
      dtype='object')
```

In [14]:

```
#Checking the categorical columns.
```

In [15]:

```
df['COUNTRY'].unique()
```

Out[15]:

```
array(['USA', 'France', 'Norway', 'Australia', 'Finland', 'Austria', 'UK',  
      'Spain', 'Sweden', 'Singapore', 'Canada', 'Japan', 'Italy',  
      'Denmark', 'Belgium', 'Philippines', 'Germany', 'Switzerland',  
      'Ireland'], dtype=object)
```

In [16]:

```
df['DEALSIZE'].unique()
```

Out[16]:

```
array(['Small', 'Medium', 'Large'], dtype=object)
```

In [17]:

```
df['PRODUCTCODE'].unique()
```

Out[17]:

```
array(['S10_1678', 'S10_1949', 'S10_2016', 'S10_4698', 'S10_4757',  
      'S10_4962', 'S12_1099', 'S12_1108', 'S12_1666', 'S12_2823',  
      'S12_3148', 'S12_3380', 'S12_3891', 'S12_3990', 'S12_4473',  
      'S12_4675', 'S18_1097', 'S18_1129', 'S18_1342', 'S18_1367',  
      'S18_1589', 'S18_1662', 'S18_1749', 'S18_1889', 'S18_1984',  
      'S18_2238', 'S18_2248', 'S18_2319', 'S18_2325', 'S18_2432',  
      'S18_2581', 'S18_2625', 'S18_2795', 'S18_2870', 'S18_2949',  
      'S18_2957', 'S18_3029', 'S18_3136', 'S18_3140', 'S18_3232',  
      'S18_3259', 'S18_3278', 'S18_3320', 'S18_3482', 'S18_3685',  
      'S18_3782', 'S18_3856', 'S18_4027', 'S18_4409', 'S18_4522',  
      'S18_4600', 'S18_4668', 'S18_4721', 'S18_4933', 'S24_1046',  
      'S24_1444', 'S24_1578', 'S24_1628', 'S24_1785', 'S24_1937',  
      'S24_2000', 'S24_2011', 'S24_2022', 'S24_2300', 'S24_2360',  
      'S24_2766', 'S24_2840', 'S24_2841', 'S24_2887', 'S24_2972',  
      'S24_3151', 'S24_3191', 'S24_3371', 'S24_3420', 'S24_3432',  
      'S24_3816', 'S24_3856', 'S24_3949', 'S24_3969', 'S24_4048',  
      'S24_4258', 'S24_4278', 'S24_4620', 'S32_1268', 'S32_1374',  
      'S32_2206', 'S32_2509', 'S32_3207', 'S32_3522', 'S32_4289',  
      'S32_4485', 'S50_1341', 'S50_1392', 'S50_1514', 'S50_4713',  
      'S700_1138', 'S700_1691', 'S700_1938', 'S700_2047', 'S700_2466',  
      'S700_2610', 'S700_2824', 'S700_2834', 'S700_3167', 'S700_3505',  
      'S700_3962', 'S700_4002', 'S72_1253', 'S72_3212'], dtype=object)
```

In [18]:

```
df.drop(['COUNTRY', 'ORDERDATE'], inplace=True, axis=1)
```

In [20]:

```
df['PRODUCTCODE'] = pd.Categorical(df['PRODUCTCODE']).codes
```

In [21]:

```
df.columns
```

Out[21]:

```
Index(['QUANTITYORDERED', 'PRICEEACH', 'ORDERLINENUMBER', 'SALES', 'QTR_ID',  
      'MONTH_ID', 'YEAR_ID', 'PRODUCTLINE', 'MSRP', 'PRODUCTCODE',  
      'DEALSIZE'],  
      dtype='object')
```

In [22]:

```
df1=pd.get_dummies(df['DEALSIZE'])
```

In [23]:

```
df1
```

Out[23]:

	Large	Medium	Small
0	0	0	1
1	0	0	1
2	0	1	0
3	0	1	0
4	0	1	0
...
2818	0	0	1
2819	0	1	0
2820	0	1	0
2821	0	0	1
2822	0	1	0

2823 rows × 3 columns

In [24]:

```
df2=pd.get_dummies(df['PRODUCTLINE'])
```

In [25]:

df2

Out[25]:

	Classic Cars	Motorcycles	Planes	Ships	Trains	Trucks and Buses	Vintage Cars
0	0	1	0	0	0	0	0
1	0	1	0	0	0	0	0
2	0	1	0	0	0	0	0
3	0	1	0	0	0	0	0
4	0	1	0	0	0	0	0
...
2818	0	0	0	1	0	0	0
2819	0	0	0	1	0	0	0
2820	0	0	0	1	0	0	0
2821	0	0	0	1	0	0	0
2822	0	0	0	1	0	0	0

2823 rows × 7 columns

In [26]:

df.drop(['PRODUCTLINE', 'DEALSIZE'], inplace=True, axis=1)

In [27]:

df=pd.concat([df,df1,df2],axis=1)

In [28]:

df.columns

Out[28]:

```
Index(['QUANTITYORDERED', 'PRICEEACH', 'ORDERLINENUMBER', 'SALES', 'QTR_ID',
      'MONTH_ID', 'YEAR_ID', 'MSRP', 'PRODUCTCODE', 'Large', 'Medium',
      'Small', 'Classic Cars', 'Motorcycles', 'Planes', 'Ships', 'Train
s',
      'Trucks and Buses', 'Vintage Cars'],
      dtype='object')
```


In [29]:

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2823 entries, 0 to 2822
Data columns (total 19 columns):
#   Column                Non-Null Count  Dtype
---  -
0   QUANTITYORDERED       2823 non-null  int64
1   PRICEEACH              2823 non-null  float64
2   ORDERLINENUMBER       2823 non-null  int64
3   SALES                  2823 non-null  float64
4   QTR_ID                2823 non-null  int64
5   MONTH_ID              2823 non-null  int64
6   YEAR_ID               2823 non-null  int64
7   MSRP                  2823 non-null  int64
8   PRODUCTCODE           2823 non-null  int8
9   Large                 2823 non-null  uint8
10  Medium                2823 non-null  uint8
11  Small                 2823 non-null  uint8
12  Classic Cars          2823 non-null  uint8
13  Motorcycles           2823 non-null  uint8
14  Planes                2823 non-null  uint8
15  Ships                 2823 non-null  uint8
16  Trains                2823 non-null  uint8
17  Trucks and Buses      2823 non-null  uint8
18  Vintage Cars          2823 non-null  uint8
dtypes: float64(2), int64(6), int8(1), uint8(10)
memory usage: 206.9 KB
```

In [30]:

```
df.head()
```

Out[30]:

	QUANTITYORDERED	PRICEEACH	ORDERLINENUMBER	SALES	QTR_ID	MONTH_ID	YE
0	30	95.70	2	2871.00	1	2	
1	34	81.35	5	2765.90	2	5	
2	41	94.74	2	3884.34	3	7	
3	45	83.26	6	3746.70	3	8	
4	49	100.00	14	5205.27	4	10	

In [31]:

```
df['PRODUCTCODE'].unique()
```

Out[31]:

```
array([ 0,  1,  2,  3,  4,  5,  6,  7,  8,  9, 10, 11, 12,
        13, 14, 15, 16, 17, 18, 19, 20, 21, 22, 23, 24, 25,
        26, 27, 28, 29, 30, 31, 32, 33, 34, 35, 36, 37, 38,
        39, 40, 41, 42, 43, 44, 45, 46, 47, 48, 49, 50, 51,
        52, 53, 54, 55, 56, 57, 58, 59, 60, 61, 62, 63, 64,
        65, 66, 67, 68, 69, 70, 71, 72, 73, 74, 75, 76, 77,
        78, 79, 80, 81, 82, 83, 84, 85, 86, 87, 88, 89, 90,
        91, 92, 93, 94, 95, 96, 97, 98, 99, 100, 101, 102, 103,
        104, 105, 106, 107, 108], dtype=int8)
```

In [32]:

```
df.columns
```

Out[32]:

```
Index(['QUANTITYORDERED', 'PRICEEACH', 'ORDERLINENUMBER', 'SALES', 'QTR_ID',
      'MONTH_ID', 'YEAR_ID', 'MSRP', 'PRODUCTCODE', 'Large', 'Medium',
      'Small', 'Classic Cars', 'Motorcycles', 'Planes', 'Ships', 'Train
      s',
      'Trucks and Buses', 'Vintage Cars'],
      dtype='object')
```

In [33]:

```
from sklearn.cluster import k_means, KMeans
from sklearn.decomposition import PCA
```


Plotting the Elbow Plot to determine the number of clusters.

In [34]:

```
distotions=[]#within clusters sum of squares from the centroid
k=range(1,10)
for i in k:
    kmeanModel=KMeans(n_clusters=i)
    kmeanModel.fit(df)
    distotions.append(kmeanModel.inertia_) #appending inertia to list
```

In [36]:

```
plt.figure(figsize=(20,10))
plt.plot(k,distotions,'bx-')
plt.xlabel('k')
plt.ylabel('Distortion')
plt.title('The elbow method')
plt.show()
```



In [37]:

```
x_train=df.values
```

In [38]:

```
x_train
```

Out[38]:

```
array([[ 30. ,  95.7 ,   2. , ...,   0. ,   0. ,   0. ],
       [ 34. ,  81.35,   5. , ...,   0. ,   0. ,   0. ],
       [ 41. ,  94.74,   2. , ...,   0. ,   0. ,   0. ],
       ...,
       [ 43. , 100. ,   4. , ...,   0. ,   0. ,   0. ],
       [ 34. ,  62.24,   1. , ...,   0. ,   0. ,   0. ],
       [ 47. ,  65.52,   9. , ...,   0. ,   0. ,   0. ]])
```

In [39]:

```
x_train.shape
```

Out[39]:

```
(2823, 19)
```

In [40]:

```
model=KMeans(n_clusters=3,random_state=2)
```

In [41]:

```
model.fit(x_train)
```

Out[41]:

```
KMeans(n_clusters=3, random_state=2)
```

In [42]:

```
pred=model.predict(x_train)
```

In [43]:

```
unique,counts=np.unique(pred,return_counts=True)
```

In [44]:

```
counts=counts.reshape(1,3)
```

In [45]:

```
counts_df=pd.DataFrame(counts,columns=['Cluster1','Cluster2','Cluster3'])
```

In [46]:

```
counts_df
```

Out[46]:

	Cluster1	Cluster2	Cluster3
0	1083	1367	373

In [47]:

```
unique
```

Out[47]:

```
array([0, 1, 2])
```

In [48]:

```
pca = PCA(n_components=2) #Converting all the features into 2 columns to make it easy to visualize using Principal Component Analysis.
```

In [50]:

```
reduced_X = pd.DataFrame(pca.fit_transform(x_train),columns=['PCA1','PCA2']) #Creating a DataFrame.
```

In [51]:

```
reduced_X.head()
```

Out[51]:


	PCA1	PCA2
0	-682.488323	-42.819535
1	-787.665502	-41.694991
2	330.732170	-26.481208
3	193.040232	-26.285766
4	1651.532874	-6.891196

In [52]:

```
#Plotting the normal Scatter Plot  
plt.figure(figsize=(14,10))  
plt.scatter(reduced_X['PCA1'],reduced_X['PCA2'])
```

Out[52]:

<matplotlib.collections.PathCollection at 0x29847b8aa00>



In [53]:

```
model.cluster_centers_ #Finding the centriods. (3 Centriods in total. Each Array contains a centroids for particular feature )
```

Out[53]:

```
array([[ 3.72031394e+01,  9.52120960e+01,  6.44967682e+00,
         4.13868425e+03,  2.72022161e+00,  7.09879963e+00,
         2.00379409e+03,  1.13248384e+02,  5.04469067e+01,
         2.08166817e-17,  1.00000000e+00, -6.66133815e-16,
         3.74884580e-01,  1.15420129e-01,  9.41828255e-02,
         8.21791320e-02,  1.84672207e-02,  1.16343490e-01,
         1.98522622e-01],
       [ 3.08302853e+01,  7.00755230e+01,  6.67300658e+00,
         2.12409474e+03,  2.71762985e+00,  7.09509876e+00,
         2.00381127e+03,  7.84784199e+01,  6.24871982e+01,
         6.93889390e-18,  6.21799561e-02,  9.37820044e-01,
         2.64813460e-01,  1.21433797e-01,  1.29480614e-01,
         1.00219459e-01,  3.87710315e-02,  9.21726408e-02,
         2.53108998e-01],
       [ 4.45871314e+01,  9.98931099e+01,  5.75603217e+00,
         7.09596863e+03,  2.71045576e+00,  7.06434316e+00,
         2.00389008e+03,  1.45823056e+02,  3.14959786e+01,
         4.20911528e-01,  5.79088472e-01,  1.66533454e-16,
         5.33512064e-01,  1.07238606e-01,  7.23860590e-02,
         2.14477212e-02,  1.07238606e-02,  1.31367292e-01,
         1.23324397e-01]])
```

In [54]:

```
reduced_centers = pca.transform(model.cluster_centers_) #Transforming the centroids into 3 in x and y coordinates
```

In [55]:

```
reduced_centers
```

Out[55]:


```
array([[ 5.84994044e+02, -4.36786931e+00],
       [-1.43005891e+03,  2.60041009e+00],
       [ 3.54247180e+03,  3.15185487e+00]])
```

In [56]:

```
plt.figure(figsize=(14,10))
plt.scatter(reduced_X['PCA1'],reduced_X['PCA2'])
plt.scatter(reduced_centers[:,0],reduced_centers[:,1],color='black',marker='x',s=300) #
Plotting the centriods
```

Out[56]:

<matplotlib.collections.PathCollection at 0x29845f81df0>



In [57]:

```
reduced_X['Clusters'] = pred #Adding the Clusters to the reduced dataframe.
```

In [58]:

```
reduced_X.head()
```

Out[58]:

	PCA1	PCA2	Clusters
0	-682.488323	-42.819535	1
1	-787.665502	-41.694991	1
2	330.732170	-26.481208	0
3	193.040232	-26.285766	0
4	1651.532874	-6.891196	0


In [59]:

```
#Plotting the clusters
plt.figure(figsize=(14,10))
# taking the cluster number and first column taking the s
ame cluster number and second column Assigning the color
plt.scatter(reduced_X[reduced_X['Clusters'] == 0].loc[:, 'PCA1'], reduced_X[reduced_X['Clusters'] == 0].loc[:, 'PCA2'], color='slateblue')
plt.scatter(reduced_X[reduced_X['Clusters'] == 1].loc[:, 'PCA1'], reduced_X[reduced_X['Clusters'] == 1].loc[:, 'PCA2'], color='springgreen')
plt.scatter(reduced_X[reduced_X['Clusters'] == 2].loc[:, 'PCA1'], reduced_X[reduced_X['Clusters'] == 2].loc[:, 'PCA2'], color='indigo')

plt.scatter(reduced_centers[:,0], reduced_centers[:,1], color='black', marker='x', s=300)
```

Out[59]:

```
<matplotlib.collections.PathCollection at 0x29833eb4580>
```



In []: